

Task 2: Values Audit

1. Mission Statement and Company Profile

The hypothetical company, referred to as FairLens RegTech, is a growth-stage specialist compliance software provider focused exclusively on AI-driven consumer credit. Its mission is to operationalise legally relevant fairness across jurisdictions — giving regulated lenders the monitoring, explanation, and governance evidence they need to satisfy both US fair lending obligations and UK FCA Consumer Duty requirements, without requiring them to build that capability in-house.

Core values: regulatory defensibility (every finding in the tool must trace to a documented regulatory source), fairness (the tool must surface real risk, not label-wash it), transparency (the rule logic, thresholds, and provenance are visible and versioned), and efficiency (the tool must reduce compliance cost, not add to it). These values are in genuine tension: tighter monitoring can reduce credit volume, increase manual review cost, and surface model weaknesses that the client would prefer not to document. The company accepts this tension as the cost of building a tool that measures substance rather than performing compliance.

The company is at Series A stage, approximately four years old. Current headcount is 45 FTE, distributed across credit analytics (12), compliance engineering (14), model governance and legal advisory (8), and sales and customer success (11). The company currently serves 12 pilot clients across the United States and United Kingdom.

Revenue aspirations:

Target ARR of USD 6-8 million within three years, derived from 50-70 mid-market lenders at an average annual contract value of USD 100,000-150,000, inclusive of implementation and ongoing support. Current ARR is approximately USD 900,000 from pilot engagements. The pricing model is subscription-based with a usage component tied to monthly active applicant volumes, creating revenue upside as clients scale.

Cost structure:

Engineering and analytics talent accounts for approximately 60% of operating expenditure. Legal and regulatory advisory retainers account for 18%, reflecting the cost of keeping rule configurations current as regulations evolve. Sales, marketing, and customer success account for 15%. Infrastructure, data, and licensing account for the remainder. Target gross margin at scale is 68-72%, consistent with comparable B2B SaaS compliance platforms. Customer acquisition cost is currently high (12-18 month payback) because enterprise compliance sales cycles are long; the target is to reduce this to 9 months by year three through partnerships with credit bureau data providers.

Geographical coverage:

The current go-to-market focus is the United States and United Kingdom. Both markets have active AI-lending sectors, established regulatory expectations, and large enough compliance teams to represent paying clients. The product roadmap targets EU expansion no earlier than 2027, aligned with the AI Act's enforcement timeline for high-risk AI systems in financial services. Singapore is a secondary priority given MAS model risk management guidance, but requires a separate sales and regulatory localisation effort.

2. Whose Perspective the Tool Serves

The primary paying client is the Chief Compliance Officer (CCO) and senior management of a regulated mid-market lender. The most materially affected stakeholder is the rejected credit applicant, who has no contractual relationship with the RegTech provider and no direct voice in tool design. The regulator is a secondary audience: it does not purchase the tool but its expectations define what the tool must evidence.

Stakeholder tensions are real. The CCO wants defensible automation and manageable compliance cost. The lender's risk management function wants predictive performance, which sometimes conflicts with fairness objectives — a model with high predictive accuracy may produce larger group disparities than a constrained fair model. Consumers want access to credit and clear explanations when they are declined. The regulator wants evidence of fair treatment and good customer outcomes, which requires the lender to document uncomfortable findings, not just clear ones. These tensions are not resolvable by tool design alone; they are political choices that the RegTech provider inherits.

Why the CCO, and have we sized the market?

The design choice to serve the CCO primarily is driven by two factors: demonstrated ability to pay and team competence. On ability to pay: Grand View Research estimated global RegTech market size at USD 21.7 billion in 2023, growing at approximately 22% compound annual growth rate, of which credit-compliance tooling is a fast-growing sub-segment. Within the US, FDIC 2024 data identifies approximately 1,200 FDIC-insured commercial banks and credit unions in the mid-market segment — defined here as institutions with USD 500 million to USD 10 billion in assets — that use automated credit scoring. Estimated annual compliance spend per institution in this segment ranges from USD 800,000 to USD 2.5 million. AI fairness and model governance tooling is an emerging but growing sub-category within that spend.

Ability to pay is further evidenced by regulatory fine exposure. US ECOA enforcement actions have resulted in settlements ranging from USD 1 million to over USD 100 million (the BancorpSouth settlement in 2016 and the Trident Mortgage settlement in 2021 are notable examples). A USD 100,000-150,000 annual compliance tool is a rational insurance purchase against that exposure. In the UK, FCA Consumer Duty enforcement is at an early stage but the FCA has made clear that board-level evidence of outcome monitoring is a prerequisite for ongoing authorisation in consumer credit. The addressable UK mid-market segment numbers approximately 120 FCA-authorized consumer credit lenders using automated scoring. On team competence: the founding team combines credit risk modelling, regulatory affairs, and software engineering experience, which is the specific combination required to translate regulatory expectations into operational software.

3. Is the Tool Measuring Genuine Risk or Documenting Compliance?

The tool is designed to measure genuine risk. One concrete design choice illustrates this commitment.

Sensitive demographic attributes — nationality, ethnicity, religion, and marital status — are excluded from the credit model's training features. This is standard practice. However, they are retained in the dataset and used in post-decision fairness auditing. This additional step is not legally required for a basic model build. A lender could claim technical compliance by simply confirming that protected attributes are absent from the model inputs. The tool goes further by testing whether model outcomes differ across groups defined by those attributes even when the attributes were not used as direct inputs. This detects the kind of indirect or proxy-driven disparity — where correlated socioeconomic variables produce group outcome differences — that would persist invisibly under a checklist approach. Including this feature costs the client more difficult conversations; excluding it would be easier to sell. It is included because it improves genuine risk detection.

A feature deliberately excluded further illustrates the design philosophy. The prototype does not produce an automated compliance certificate or a single pass/fail score. Such a feature would be commercially attractive and easy to implement. It would also serve documentation over substance: a single score encourages clients to treat the output as an audit pass rather than as evidence requiring human review. The tool instead produces specific findings linked to specific rule configurations and source references, and explicitly flags which findings require human judgement rather than automated resolution.

4. Who Bears the Cost if the Tool Gets It Wrong?

False positives — flagging a compliant portfolio as non-compliant — impose cost on the lender through unnecessary manual review and investigation. More seriously, a lender may respond to repeated false positives by tightening credit criteria to clear the tool's thresholds. This can reduce lending access to the very groups the tool was designed to protect. The stakeholder who ultimately bears this cost is the consumer who is denied credit unnecessarily because the RegTech tool overcalibrated.

False negatives — clearing a non-compliant portfolio — are more serious in terms of harm magnitude and harder to detect. A lender that believes its model has cleared a fairness review may reduce monitoring frequency and defer remediation. If hidden disparity persists undetected, affected consumer groups face systematic credit exclusion without any regulatory response or internal correction. The harm is diffuse, cumulative, and invisible to the individuals affected. The specific stakeholder who bears this cost is the consumer applicant in the disadvantaged group.

Performance drift imposes cost on both lenders and consumers when economic conditions or applicant mix shift without detection. A model that was validated as fair under one macroeconomic environment may produce materially different disparity ratios two years later when income distributions or employment patterns change. Without continuous drift monitoring, neither the lender nor the regulator has visibility into this degradation. The prototype addresses this with a recession-scenario drift simulation, but real drift monitoring requires ongoing data ingestion that the prototype does not provide.

Jurisdiction misconfiguration — applying the US rule configuration to a UK-regulated product, or vice versa — may expose the lender to regulatory action in the incorrectly configured market. If a UK-regulated entity's product is accidentally evaluated under the less demanding US configuration, the FCA Consumer Duty evidence requirements will not be checked, and the board report will not reflect the correct risk picture. The lender bears this cost in the form of regulatory exposure; the consumer bears it in the form of reduced protection.

The most dangerous failure is a false sense of compliance: the tool produces thorough-looking reports and findings while genuine harmful outcomes persist in the portfolio. This failure is the hardest to detect, because it requires the tool to be honest about the limits of its own coverage — which is commercially inconvenient. For this reason, the tool explicitly documents what it does not check, what data it does not have access to, and where human judgement is required rather than automated. The boundary is drawn deliberately to prevent the tool from being used as a shield against regulatory scrutiny rather than as a genuine risk instrument.